

How we designed **BeThere**

A human-centred design portfolio: from a 43-person survey and three rounds of interviews to a shipped app that turns a loose "maybe" into a committed plan, with no organiser and no public maybe.

survey n=43


3 interview rounds

9 named participants

live builds, M2-M4


How we ran the research

We worked the **Double Diamond** across the term, pairing each phase to a milestone. From M3 the build itself was the probe: people drove a deployed product on their own device, so every later finding came from a real person meeting a real artefact.



METHOD	WHO AND WHEN	WHY THIS METHOD
<i>Friend Meetup Dynamics survey</i>	43 respondents, 19-20 May 2026	Tells you how widespread a behaviour is - whether plans dying is common, and what people do when unsure - before committing to a build.
<i>Round 1 discovery interviews</i>	Luca, Luke, Matthew, Noah's friend; pre-build	A survey cannot tell you why someone sends a "maybe". Semi-structured interviews surfaced the felt mechanism behind the numbers.
<i>Round 2 iteration usability</i>	Tom, Felicity, Nathan, Fiza; M3, think-aloud on the live app	Reacting to a working artefact, not a description, is where the sharpest findings came: anonymity validated, the conditional challenged.
<i>Round 3 M4 usability</i>	Luca, Nathan, Zack Foreman, Will Groves, Fangyi Lin, Thomas Gonzalez; M4, live web funnel	Testing onboarding, web-link joining and plus-ones on the real funnel, with people who had not seen the app.

THE RISKIEST ASSUMPTION WE TESTED EACH ROUND



REFLECTION

Prototype-led research was our single most useful technique. As Lukas put it on the design board, "if we start with an implementation in code then we can use that to dig into people's brains." Deploying each interaction model and watching people use it produced sharper findings than asking ever did - the hard part was staying non-leading, which is why we coached each other to drop the jargon.

The people we designed for

When only slightly unsure they can make a plan, **37 of 43 (86%)** send a "maybe" or nothing at all. Three of our four personas are defined by how they avoid committing; the fourth is the minority BeThere wants to make ordinary.

 photo RELUCTANT INITIATOR Luca 4/43 "mostly me" "I'm normally the only one who actually wants to do it." Float an idea without it being <i>his</i> idea - anonymity.	 photo THE HEDGER Matthew 27/43 send "a maybe" "Maybe means I keep the time semi-free... then I ditch." Remove "maybe": force yes / no / conditional, plus a deadline.
 photo THE GHOST Noah's friend 10/43 send nothing "It's just people, ignore the message." A private +1 that only surfaces if it matches - silence is no longer a dead end.	 photo THE DECISIVE ONE Tom the rare clean responder "Make it anonymous and there's no pressure on that person at all." A deadline and push - plus the anonymity that makes everyone answer.

A tension we kept honest (Felicity, R2): naming who you will "go if" feels exposing even when private. We carried that split forward rather than designing it away.

Who the product touches

BeThere is not competing with "no tool" - it competes with the silent **WhatsApp thread**, and leans on hosted services. Stakeholders sit in rings by how directly the product touches them.

INDIRECT - AFFECTED, NOT USERS

The group chat / WhatsApp (the incumbent) · Venues and hosts · Hosted services (Clerk, AWS, Vercel)

DIRECT - DRAWN INTO A PLAN, NOT DRIVING IT

Wider group members · People named in a conditional · Plus-ones and cross-group guests

CORE - THE PEOPLE IT IS FOR

The Reluctant Initiator · The Decisive One · The Hedger · The Ghost

CORE NEED

Move a plan without bearing the social cost of being the pusher. 4/43 lead; Tom, Felicity, Luca.

DIRECT NEED

Show interest without the maybe-cascade, and never be named awkwardly. 27/43 maybe; Felicity's tension.

INDIRECT NEED

Beat WhatsApp on convenience, or users stay put. Felicity: "why bother if you're on WhatsApp anyway?"

The hardest stakeholder is the one we replace: Felicity's "a website would be less stressful" is why BeThere later shipped a web target and a no-download share link (DRP-56). As Luke put it, when a plan dies in the thread, "no one says shit".

What the research told us

After dropping a navigation app ("Stumble") and a "prediction market for your friend group", we converged on meetup coordination around **19 May**. The seed was Noah's discovery note: a friend whose meetups "fail about once a week... people not saying whether or not they can make it" - "if only there was a discrete way to say whether you can make it".

23/43

said 30% or more of recent hangouts were talked about but never happened.

36/43

say smaller groups make meeting easier; none said larger.

20/43

would use a private "up for something" signal more than messaging the group.

10/43

say the first suggested time gets agreed often - so BeThere votes across candidates.

VALIDATED INSIGHTS (SURVEY + INTERVIEWS)

- 01 The bottleneck is suggestion and commitment, not execution.** Plans die before a time is fixed; people who show up are fine. (Luca; survey: first time rarely agreed.)
- 02 The "maybe" and silence are the failure mechanism, and uncertainty is contagious.** Luke: "it cascades uncertainty on the group... that is a website written off". Survey: 86% reply maybe or nothing.
- 03 Effort is asymmetric; nobody wants to be the organiser.** Only 4/43 mostly lead. (Felicity, Tom, Luca.)
- 04 Anonymity removes the social cost of initiating** - a plan feels like consensus, not one person's ask. (Tom; 20/43 want a private signal.)

IN THEIR OWN WORDS - WHY PLANS DIED

"Couldnt be arsed." · "Too tired to go potentially because its too far." · "Something else came up that I wanted to go to more than the meet up." · "My failure to plan accordingly, eg I forgot."

Low-stakes drift - the kind a blind, deadlined commitment is designed to hold together.

THE OPPORTUNITY

How might we help a group of friends turn a loose intention to meet into a committed plan, without anyone having to be the organiser, and without the social pressure of a public "maybe"?

Journey maps: before and after

The plan does not die from a logistics failure; it dies from a feelings failure - so both journeys track **emotion** as the primary signal. The before-curve only ever falls; the future-state's job is not to add features but to stop the fall.

BEFORE · in a dead chat

STAGE	EMOTION
1. The spark "We should do something soon!!"	Hopeful
2. The hedge A few "maybe!", two go silent.	Guarded
3. The cascade The first "I can't" lands; replies thin.	Deflated
4. The drag Days pass; nobody re-raises it.	Apathy
5. The doubt creep Second thoughts settle; excuses surface.	Regret
6. Silent death No message ends it; the thread stops.	Disappointed

AFTER · the same friends, on BeThere

STAGE	EMOTION
1. Send (anonymous) One meetup: candidate times, a deadline.	Relieved
2. Collect and vote Members +1; counts show, names never.	Safe
3. Momentum, not maybe No "maybe" button, only a +1 or no vote.	Encouraged
4. Auto-lock Top time wins at the deadline; no one calls it.	Fair
5. The blind moment Private RSVP: Yes / Can't / "I'll go if".	Decisive
6. It clears or fizzles "You and 4 others are in", or silence.	Confident

"Days went past... it just got more drawn out... social thoughts ended up creeping in and more and more people ended up pulling out."

- Matthew, R1

"When someone expresses some kind of uncertainty, it cascades uncertainty on the group... that is a website written off."

- Luke, R1

Service blueprint

The future journey mapped onto the real system: frontstage actions, the mobile screens, the backstage tRPC procedures and the Postgres data. There is **no scheduler** - phase transitions settle lazily on every read or write, so a plan auto-locks and clears or fizzles the next time anyone touches it, keeping the privacy boundary inside the server.

FRONTSTAGE	Send one meetup (anonymous)	+1 times / activities; add a candidate	Deadline passes (no user action)	RSVP: Yes / Can't / "I'll go if [people]"	See "You're in" + the in-crowd, or nothing
MOBILE SCREENS	CreateWizard	EventDetail (collecting: public counts, no names)	Dashboard (Going / Open / Done)	EventDetail (blind moment)	EventDetail (cleared)
BACKSTAGE TRPC + LOGIC	events.create	toggleReaction / addCandidate / setOptOut	settlePhase -> openMoment, pickWinnerOrBestId	events.respond (conditional server-side)	settlePhase : clears(response, quorum)
POSTGRES	events	event_candidates, candidate_reactions	events	responses (kind, cond)	events

PRIVACY BOUNDARY - GUARDS EVERY READ

isAnonymous defaults true and the read layer returns **isCreator** as a private boolean only, never the creator's id.

No live tally. Collecting returns per-candidate +1 totals but never the voter set; the moment returns nothing about others.

Silent fizzle. A plan that misses quorum flips to **fizzled** on the next read and drops out of every dashboard - no push, no dead thread.

Three privacy promises are structural, not copy: they live in the server's procedures, not in the words on the screen. Names from `events.ts` and `db/schema.ts`.

Prototyping: the fidelity ladder

Prototyping was a deliberate **fidelity ladder**, raised one rung only when a real question ran out of answers. Paper showed the journey made sense but not whether the interaction model was right; the skeleton gave behaviour but we drove it ourselves; live builds let a friend use the actual product on their own device, which is what made the iteration interviews **genuine usability tests**.

R1	HAND-DRAWN M2 MOCK-UPS DRP-17 Is the screen-to-screen flow legible?
R2	CLICKABLE WALKING SKELETON DRP-14/15 Does the whole loop hold together end to end? Real Postgres, full loop.
R3	PUBLICLY DEPLOYED LIVE BUILDS DRP-16/24/31 Will a real friend, on their own phone, get it? Phone + browser, real Clerk auth.
R4	SUCCESSIVE CONCEPT PROTOTYPES DRP-20/29/30/41 Which interaction model do users actually want? Concrete RSVP, convergence, float, unified suggest flow.

RUNG 1 - AUTHENTIC ARTEFACTS

Hand-drawn create screen + "I'll make it if..." moment sheet
(docs/mockups/m2/ALL MOCKUPS.pdf)

The create sketch became **CreateWizard**; the moment sheet with friends and tick boxes is the direct ancestor of the `events.respond` conditional.

RUNG 3 - LIVE-APP SCREENS

Create wizard

Collecting: public counts, no names

Cleared "You're in" moment

Three screens from the live build, showing the refined-neobrutalist visual system shipped in DRP-25.

From finding to change

The artefact that kept us honest: every change we shipped points back at a **real survey number or interview quote**, and forward at a **real DRP issue**. If a finding did not exist, the feature did not get built.

FINDING (SOURCE)	INSIGHT	CHANGE SHIPPED
27/43 send "a maybe", 10/43 go silent when slightly unsure; only 5 give a clean answer. (<i>survey</i>)	Hedging is the default, and the failure mechanism.	A blind, deadlined RSVP - Yes / Can't make it / "I'll go if [people]" - no public maybe, no running tally. DRP-29 / 41
"The organisers should have an option to be anonymous... no pressure on that person at all" (<i>Tom</i>); only 4/43 mostly lead.	Anonymity reframes a plan as group consensus.	Creator anonymity is always on - the group sees the meetup, never who proposed it. No toggle. DRP-41
"You split the load across your friends... without feeling like you're dragging everyone" (<i>Felicity</i>); 36/43 say bigger groups are harder.	People want to co-decide, not be organised.	Two member-fed candidate lists, TIME and ACTIVITY, anyone can add to and +1; counts public, names never. DRP-30 / 41
"A checkbox or a switch... this is a concrete event, you cannot change it... like a football match" (<i>Tom</i>); confusion over three suggest modes.	One switch should cover concrete and open, not three modes.	lockTimes / lockActivity flags; a fully pinned plan skips straight to a blind moment. The three-mode fork was deleted. DRP-41
"The deadline would be better... but only if there was a push... like BeReal, it goes off, ding ding ding" (<i>Tom</i>).	A blind, deadlined, push-driven yes/no beats an open poll.	Lock-in deadline plus device-local "Voting closes" / "RSVP closes" reminders. DRP-32 / 41
The countdown bar turned greener as the deadline neared, signalling "you're nearly there" exactly when urgency should rise. (<i>M3 review</i>)	Urgency was being signalled backwards.	The green drain bar became a plain-language countdown ("Voting closes in 2 days"); urgency shown in pink under one hour. DRP-43
Felicity could not name a friend in a conditional, even privately (<i>R2</i>); Tom liked naming people.	Felt privacy is not the same as technical privacy.	The "Go if [people]" copy was softened and defaulted to "At least one" rather than a named person. DRP-47
"Hoping people download it is probably not it" (<i>Fangyi Lin</i>); "a website would be less stressful" (<i>Felicity</i>); "it's just a link" (<i>Luca, R3</i>).	The download is the adoption barrier, not the idea.	Invite links and join codes, plus a no-download meetup share link that converts to a responding member. DRP-50 / 56

Impact

We frame impact modestly: a 43-response survey, three interview rounds and a working build, but **no longitudinal usage data**. The claims below are the change the design is *intended* to produce, told through the personas and tied to evidence that the underlying problem is real.

DIRECT IMPACT, THROUGH THE PERSONAS

Luca Reluctant Initiator

Floats a plan without it being his ask.
Anonymity (DRP-41) means no one to chase and no one to blame - the asymmetric effort he lives with: *"I'm normally the only one who actually wants to do it."*

Matthew The Hedger

No maybe to hide behind and no infinite runway. The deadline (DRP-32/41) closes the plan before his "second thoughts" can grow.

The Ghost Noah's friend

A quiet +1 is now a safe, useful signal, not a dead end. Names are never shown; the plan fires only if enough others matched him.

Tom The Decisive One

The blind, deadlined moment with a push gives his clean yes/no a system that makes everyone else answer too.

WIDER IMPACT, CAREFULLY BOUNDED

Term time disperses groups, and the coordination cost of regrouping is the cost BeThere lowers. Matthew described a reunion that "got more drawn out... and more and more people ended up pulling out", and the survey shows the failure is common. Venues and hosts get a firm time and headcount the moment a plan clears; the group chat is augmented, not replaced. We do not over-claim: BeThere does not cure loneliness, and we have not measured wellbeing. The honest statement is narrower - **it removes one evidence-backed friction** that causes low-stakes plans to die.

Impact asset: "The Gathering"
speculative 2028 cover story - no
organiser, no maybe, no names. Every
stat on it is real DRP_02 primary
research.

The cover story imagines that future explicitly: a 2028 magazine cover framing the headline outcomes - no organiser, no public maybe, shared labour - as forward-looking, not measured.

"I care less about what we do, I just want to see them." - Felicity

What HCD did to the product

The strongest evidence that we did human-centred design, not feature-led design, is that the product model itself changed in response to users, repeatedly. The product is **smaller and sharper** than the one we first imagined, and that narrowing is the clearest signal we were listening.

DELETED

A feature: the three-mode suggest fork collapsed into two lock flags.

DRP-41

PROMOTED

A flag to a principle: anonymity became always-on, the only mode.

DRP-41

LEFT OPEN

A real tension: naming people in a conditional, surfaced not designed away.

DRP-47

WHAT WE WOULD DO DIFFERENTLY

Prioritise onboarding earlier, so the iteration interviews tested the whole journey, not just the core loop.
Onboarding only landed in M4 (DRP-50).

OPEN WORK, FRAMED BY USER NEEDS

- The felt-privacy problem in the conditional (test Felicity's "rank importance" alternative).
- Reliable remote push (currently device-local, blocked on a dev build rather than Expo Go).
- Competing with the group chat people already live in (the web-first link, DRP-56, is the first move).

We deliberately did not collect SUS, task-success or telemetry; the quantitative-evaluation report scopes those instruments, and we will not present an intended outcome as a measured one until they are run.

REFLECTION

The project's own history is the best evidence of the process: a cancelled feature (DRP-40, scaling moment length by distance), a betting-market idea dropped before it began, three suggest modes collapsed into two flags, and a conditional softened after one participant could not face it. These are the marks of a design that kept changing because real people kept telling us to.