


How we designed **BeThere**

A human-centred design portfolio: from a 47-person survey and three rounds of interviews to a shipped app that turns a loose "maybe" into a committed plan, with no organiser and no public maybe.


How we ran the research

We worked the **Double Diamond** across the term, pairing each phase to a milestone. From M3 the build itself was the probe: people drove a deployed product on their own device, so every later finding came from a real person meeting a real artefact.



METHOD	WHO AND WHEN	WHY THIS METHOD
<i>Friend Meetup Dynamics survey</i>	47 respondents, 19-20 May 2026	Shows how widespread the behaviour is - is plan-death common, what do people do when unsure - before we commit to a build.
<i>Round 1 discovery interviews</i>	Luca, Luke, Matthew, Noah's friend; pre-build	Surfaces <i>why</i> someone sends a "maybe" - the felt mechanism behind the numbers.
<i>Round 2 iteration usability</i>	Tom, Felicity, Nathan, Fiza; M3, think-aloud on the live app	Reacting to a working artefact gave the sharpest findings: anonymity validated, the conditional challenged.
<i>Round 3 M4 usability</i>	Luca, Nathan, Zack Foreman, Will Groves, Fangyi Lin, Thomas Gonzalez; M4, live web funnel	Onboarding, web-link joining and plus-ones, tested on the real funnel with fresh users.

THE RISKIEST ASSUMPTION WE TESTED EACH ROUND



REFLECTION

Prototype-led research was our most useful technique. As Lukas put it, **"if we start with an implementation in code then we can use that to dig into people's brains."** Deploying each model and watching people use it beat asking - the hard part was staying non-leading, so we coached each other to drop the jargon.

The people we designed for

Vasanth and Milly are composites of the students we interviewed - the two sides of every failed plan: **the organiser** who burns out chasing replies, and **the participant** whose "maybe" is really uncertainty.



Vasanth

The Organiser

21 · Law, St Andrews · volleyball, board games, cooks for friends · extroverted, reliable

"I'm normally the only one who actually wants to do it."

Loves seeing his friends, but organising is **difficult** and chasing replies from hesitant members is **frustrating**. BeThere lets him float a plan anonymously, so it stops being his job to push.



Milly

The Participant

19 · English Literature, St Andrews · cello, reading · new to the netball society · shy, introverted



"I'm down if others are down."

Would love to go to more socials, but she will not commit unless she knows a friend is going too - so her reply is a hedging "maybe". BeThere lets her say yes **privately and conditionally**, with names never shown.

Grounded in the interviews: organisers (Luca, Eddie) and participants (Luke: "I'm down if others are down"). One tension we kept honest - Felicity could not name who she would "go if" even privately - so we softened that to "at least one" rather than a named person.

Who the product touches

BeThere is not competing with "no tool" - it competes with the silent **WhatsApp thread**, and leans on hosted services. Stakeholders sit in rings by how directly the product touches them.



The hardest stakeholder is the one we replace: Felicity's "a website would be less stressful" is why BeThere later shipped a web target and a no-download share link (DRP-56). As Luke put it, when a plan dies in the thread, "no one says shit".

What the research told us

After dropping a navigation app ("Stumble") and a "prediction market for your friend group", we converged on meetup coordination around **19 May**. The seed was Noah's discovery note: a friend whose meetups "fail about once a week... people not saying whether or not they can make it" - "if only there was a discrete way to say whether you can make it".

87%

send a "maybe" or delay their reply to a plan they are unsure about.

51%

say over a third of their meet-ups never make it out of the chat.

60%

say the meet-up time gets moved half the time or more.

83%

say smaller groups make it easier; almost none say larger.

VALIDATED INSIGHTS (SURVEY + INTERVIEWS)

01 The bottleneck is suggestion and commitment, not execution. Plans die before a time is fixed; people who show up are fine. (Luca; 60% say the time gets moved half the time or more.)

02 The "maybe" and silence are the failure mechanism, and uncertainty is contagious. Luke: "it cascades uncertainty on the group... that is a website written off". Survey: 87% send a maybe or delay.

03 Effort is asymmetric; nobody wants to be the organiser. Only 6 of 47 mostly lead. (Felicity, Tom, Luca.)

IN THEIR OWN WORDS - WHY PLANS DIED

"Couldnt be arsed." · "Too tired to go potentially because its too far." · "Something else came up that I wanted to go to more than the meet up." · "My failure to plan accordingly, eg I forgot."

Low-stakes drift - the kind a blind, deadlined commitment is designed to hold together.


THE OPPORTUNITY

How might we help a group of friends turn a loose intention to meet into a committed plan, without anyone having to be the organiser, and without the social pressure of a public "maybe"?


Journey maps: before and after

We mapped one real pub quiz two ways - as Vasanth (organiser) felt it, and as Milly (participant) felt it. Both curves only ever fall, for opposite reasons: **his is effort**, **hers is uncertainty**. On BeThere, the same plan inverts.


VASANTH'S PUB QUIZ, AS HE FELT IT



MILLY'S PUB QUIZ, AS SHE FELT IT




VASANTH'S NEW JOURNEY ON BETHERE (THE FAINT LINE IS THE GROUP CHAT - IT STILL DIES)



Service blueprint

The future journey runs on one typed system. There is **no scheduler** - phases settle lazily on read - and the privacy promises that make the moment feel safe are enforced in the server, out of the client's reach.



PRIVACY BOUNDARY - THE EXPERIENCE GUARANTEE, ENFORCED SERVER-SIDE

isAnonymous defaults true and the read layer returns **isCreator** as a private boolean only, never the creator's id.

No live tally. Collecting returns per-candidate +1 totals but never the voter set; the moment returns nothing about others.

Silent fizzle. A plan that misses quorum flips to **fizzled** on the next read and drops out of every dashboard - no push, no dead thread.

The anonymity and blind moment the users asked for are structural, not copy: they live in the server's procedures (`events.ts`, `db/schema.ts`), so they cannot leak to the client.

Prototyping: the fidelity ladder

Prototyping was a deliberate **fidelity ladder**, raised one rung only when a real question ran out of answers. Paper showed the journey made sense but not whether the interaction model was right; the skeleton gave behaviour but we drove it ourselves; live builds let a friend use the actual product on their own device, which is what made the iteration interviews **genuine usability tests**.

R1	HAND-DRAWN M2 MOCK-UPS DRP-17 Is the screen-to-screen flow legible?
R2	CLICKABLE WALKING SKELETON DRP-14/15 Does the whole loop hold together end to end? Real Postgres, full loop.
R3	PUBLICLY DEPLOYED LIVE BUILDS DRP-16/24/31 Will a real friend, on their own phone, get it? Phone + browser, real Clerk auth.
R4	SUCCESSIVE CONCEPT PROTOTYPES DRP-20/29/30/41 Which interaction model do users actually want? Concrete RSVP, convergence, float, unified suggest flow.

RUNG 1 - AUTHENTIC ARTEFACTS

Hand-drawn create screen + "I'll make it if..." moment sheet
(docs/mockups/m2/ALL MOCKUPS.pdf)

The create sketch became **CreateWizard**; the moment sheet with friends and tick boxes is the direct ancestor of the `events.respond` conditional.

RUNG 3 - LIVE-APP SCREENS

Create wizard

Collecting: public counts, no names

Cleared "You're in" moment

Three screens from the live build, showing the refined-neobrutalist visual system shipped in DRP-25.

From finding to change

The artefact that kept us honest: every change we shipped points back at a **real survey number or interview quote**, and forward at a **real DRP issue**. If a finding did not exist, the feature did not get built.

FINDING (SOURCE)	INSIGHT	CHANGE SHIPPED
87% send "a maybe" or delay when slightly unsure; only 6 of 47 give a clean answer. (<i>survey, n=47</i>)	Hedging is the default, and the failure mechanism.	A blind, deadlined RSVP - Yes / Can't make it / "I'll go if [people]" - no public maybe, no running tally. DRP-29 / 41
"The organisers should have an option to be anonymous... no pressure on that person at all" (<i>Tom</i>); only 6 of 47 mostly lead.	Anonymity reframes a plan as group consensus.	Creator anonymity is always on - the group sees the meetup, never who proposed it. No toggle. DRP-41
"You split the load across your friends... without feeling like you're dragging everyone" (<i>Felicity</i>); 83% say smaller groups are easier.	People want to co-decide, not be organised.	Two member-fed candidate lists, TIME and ACTIVITY, anyone can add to and +1; counts public, names never. DRP-30 / 41
"A checkbox or a switch... this is a concrete event, you cannot change it... like a football match" (<i>Tom</i>); confusion over three suggest modes.	One switch should cover concrete and open, not three modes.	lockTimes / lockActivity flags; a fully pinned plan skips straight to a blind moment. The three-mode fork was deleted. DRP-41
"The deadline would be better... but only if there was a push... like BeReal, it goes off, ding ding ding" (<i>Tom</i>).	A blind, deadlined, push-driven yes/no beats an open poll.	Lock-in deadline plus device-local "Voting closes" / "RSVP closes" reminders. DRP-32 / 41
The countdown bar turned greener as the deadline neared, signalling "you're nearly there" exactly when urgency should rise. (<i>M3 review</i>)	Urgency was being signalled backwards.	The green drain bar became a plain-language countdown ("Voting closes in 2 days"); urgency shown in pink under one hour. DRP-43
Felicity could not name a friend in a conditional, even privately (<i>R2</i>); Tom liked naming people.	Felt privacy is not the same as technical privacy.	The "Go if [people]" copy was softened and defaulted to "At least one" rather than a named person. DRP-47
"Hoping people download it is probably not it" (<i>Fangyi Lin</i>); "a website would be less stressful" (<i>Felicity</i>); "it's just a link" (<i>Luca, R3</i>).	The download is the adoption barrier, not the idea.	Invite links and join codes, plus a no-download meetup share link that converts to a responding member. DRP-50 / 56

Impact

We frame impact modestly: a 47-response survey, three interview rounds and a working build, but **no longitudinal usage data**. The claims below are the change the design is *intended* to produce, told through Vasanth and Milly and tied to evidence the problem is real.

DIRECT IMPACT, THROUGH THE PERSONAS

Vasanth The Organiser

Floats a plan anonymously, so it stops being his job to push. No chasing, no one to blame - the asymmetric effort he lives with disappears: *"I'm normally the only one who actually wants to do it."*

Milly The Participant

Says yes privately and conditionally - "I'll go if" - with names never shown. Her hedging "maybe" becomes a safe, real signal that only surfaces if enough others matched her.

WIDER IMPACT, CAREFULLY BOUNDED

The friction we ease is not trivial. **15-24s spend ~70% less time with friends in person**

than two decades ago (US Surgeon General, 2023); 40% of UK 16-24s feel lonely often or very often, the loneliest age group (BBC Loneliness Experiment, 2018); and over a quarter of lonely young people have sought therapy. Fewer dead plans ripples into university counselling load, NHS mental-health provision, society engagement and local venues. We do not claim BeThere cures loneliness - we have not measured wellbeing - but every plan that does not die is a small push against that drift.

Impact asset: "The Gathering" - the speculative 2028 cover story (no organiser, no maybe, no names). Every stat on it is real DRP_02 primary research.

Our impact asset is the cover story: a 2028 magazine cover imagining the wider future BeThere could create - shared labour, no organiser, no public maybe - as forward-looking, not measured. It folds in the M4 review feedback and is deliberately not the pitch leaflet.

"I care less about what we do, I just want to see them." - Felicity

What HCD did to the product

The strongest evidence that we did human-centred design, not feature-led design, is that the product model itself changed in response to users, repeatedly. The product is **smaller and sharper** than the one we first imagined, and that narrowing is the clearest signal we were listening.

DELETED

A feature: the three-mode suggest fork collapsed into two lock flags.

DRP-41

PROMOTED

A flag to a principle: anonymity became always-on, the only mode.

DRP-41

LEFT OPEN

A real tension: naming people in a conditional, surfaced not designed away.

DRP-47

WHAT WE WOULD DO DIFFERENTLY

Prioritise onboarding earlier, so the iteration interviews tested the whole journey, not just the core loop. Onboarding only landed in M4 (DRP-50).

OPEN WORK, FRAMED BY USER NEEDS

- The felt-privacy problem in the conditional (test Felicity's "rank importance" alternative).
- Reliable remote push (currently device-local, blocked on a dev build rather than Expo Go).
- Competing with the group chat people already live in (the web-first link, DRP-56, is the first move).

We deliberately did not collect SUS, task-success or telemetry; the quantitative-evaluation report scopes those instruments, and we will not present an intended outcome as a measured one until they are run.

REFLECTION

The project's own history is the best evidence of the process: a cancelled feature (DRP-40, scaling moment length by distance), a betting-market idea dropped before it began, three suggest modes collapsed into two flags, and a conditional softened after one participant could not face it. These are the marks of a design that kept changing because real people kept telling us to.